



Sadir Ahmed Zidan

Email: sazidan559@gmail.com | Phone: 01640423606

LinkedIn: [Sadir Ahmed Zidan](#) | Github: <https://github.com/zidan010>

Address: Nikunja-2, Dhaka, Bangladesh

AI/ML Engineer with 2+ years of production experience building and deploying LLM-powered applications. Specialized in fine-tuning LLMs with LoRA/QLoRA, designing advanced RAG & agentic systems, and delivering low-latency inference APIs using LangChain, Groq, FastAPI, Docker, and MLflow. Hands-on expertise in computer vision, 3D reconstruction and multilingual retrieval pipelines (FAISS, Milvus), along with voice enabled systems. Proven track record of shipping scalable AI solutions end-to-end, from custom model training and pretrained model integration to production deployment and monitoring.

Education

American International University Bangladesh

BSc in Computer Science & Engineering, January 2020 – September 2023

CGPA: 3.77/4.00

Work Experience

Jr. Machine Learning Engineer - Fulltime, Onsite

Nasir Syntax Solution Limited (concern of Nasir Group), January 2024 – Present

- Designed and shipped production RAG and agentic LLM systems using LangChain, FAISS, Milvus, and custom retrieval pipelines for personalized search and chatbots.
- Fine-tuned and quantized LLMs with LoRA/QLoRA, built dynamic model loading and inference optimization workflows.
- Developed computer vision pipelines including object detection, segmentation and 3D reconstruction projects.
- Trained custom deep learning models from scratch and integrated pretrained models.
- Built and deployed scalable FastAPI inference services with Docker, MLflow tracking, request validation, rate-limiting, and monitoring.

AI Developer - Part time, Remote

GradMate.ai, Apr 2025 – Nov 2025

- Built high-performance agentic RAG systems and multi-step AI workflows using LangChain and Groq's ultra-fast LLM inference.
- Implemented real-time speech-to-text (STT) and text-to-speech (TTS) integrations into educational AI agents.
- Designed and maintained production-grade FastAPI backends for LLM inference, prompt versioning, and dynamic retrieval.
- Optimized latency and cost for Groq/LangChain pipelines serving live user-facing features.
- Collaborated on modular AI architecture, data contracts, and retrieval quality standards with frontend and platform teams.

AI Engineer - Contractual, Remote

Ambitionix, May 2025 – Present

- Fine-tuned LLMs for specific tasks.
- Created APIs for deployment.
- Collaborated with cross-functional teams to achieve project goals.

Machine Learning Intern

Nasir Syntax Solution Limited (concern of Nasir Group), October 2023 – December 2023

- Conducted research on personalized recommendation systems, chatbots, and text generation using LLMs and computer vision.
- Scraped and preprocessed data, including tokenization and vectorization, and fine-tuned LLMs for specific applications.

Skills

- **Programming Languages:** Python, R
- **Frameworks & Libraries:** TensorFlow, Keras, LangChain, Groq, PyTorch, Fast API, Pandas, Faiss, Llama cpp, ONNX, Ollama, Hugging Face, Streamlit
- **Tools & Platforms:** Docker, MLFlow, Kubernetes, GitHub, VS Code, R Studio, Colab, LoRa
- **Databases:** SQL, MongoDB
- **Operating Systems:** Windows, Linux

Personal Projects

- **Multilingual RAG System for Bengali and English Queries**
 - A system to process queries in English and Bengali, retrieve relevant information from a PDF document (specifically the HSC26 Bangla 1st Paper book), and generate meaningful answers.
- **Face Mask Detection**
 - A system to detect facial masks using a cascade classifier. Used a large dataset that varies with different types of poses with masks worn by humans. The model got good accuracy after implementation.
- **Drowsiness Detection**
 - A real-time system using machine learning algorithms and image processing to detect driver drowsiness. It can continuously monitor the driver while they are driving and provide immediate feedback or alerts.
- **Sentiment Analysis Chatbot**
 - This project demonstrates a simple sentiment analysis chatbot that analyzes the user's mood based on their input and provides contextual responses.
- **Life Insurance Support Assistant**
 - Agentic life-insurance assistant built using LangChain + Groq, designed to answer domain-specific questions about policy types, benefits, eligibility, and claims, using structured category-based reasoning.
- **Voice-Enabled Agentic RAG**
 - A modular, local Retrieval-Augmented Generation (RAG) pipeline for multi-domain document Q&A, featuring PDF cleaning, adaptive chunking, FAISS vector indexing, multi-agent retrieval, and both text and voice CLI interfaces.